

Stein Variation Gradient Descent

Report on Stein Variation Gradient Descent as developed and presented in a paper by Q.Liu & D.Wang, 2019: Stein variational gradient descent: A general purpose bayesian inference algorithm

James, Zoryk

January 26, 2024

This report introduces and explores the Stein Variational Gradient Descent (SVGD) algorithm as a powerful and general-purpose tool for Bayesian Inference. Unlike conventional methods such as Markov Chain Monte Carlo (MCMC) [1, 2], SVGD leverages variational inference [3] through deterministic optimization, providing efficiency and scalability for handling large datasets. The algorithm is rooted in the theoretical connection between the Kernelized Stein Discrepancy (KSD) [4] and the Kullback-Leibler (KL) divergence, making it applicable beyond variational inference. This report outlines the algorithm, delves into its theoretical foundations, and offers insights into its practical implementation.

Background

Kernels and Feature spaces

In SVGD, kernels and feature spaces play a key role in facilitating efficient computation with high-dimensional data. Let $X \subseteq \mathbb{R}^d$ denote the input space. A function $k : X \times X \rightarrow \mathbb{R}$ is considered a kernel function if there exists an inner product space $\langle F, \langle \cdot, \cdot \rangle \rangle$ and a function $\varphi : X \mapsto F$ such that

$$k(x, x') = \langle \varphi(x), \varphi(x') \rangle$$

holds for all $x, x' \in X$. This concept is employed to find linear solutions for problems in the feature space F that may not be solvable through linearity in X . In many instances, the function φ is unknown and unnecessary, although theorems exist asserting its existence [5, 6]. In simpler terms, the kernel function k can be interpreted as a measure of similarity between two points. Consider the Gaussian Radial Basis Function (RBF) kernel as an example, defined by

$$k(x, x') = \exp(-1/h \|x - x'\|_2^2) \quad (1)$$

for a suitable choice of $h \in \mathbb{R}$, known as the bandwidth or scale parameter. The transition to the higher-dimensional function space becomes evident by examining the classic example depicted in the code 1 with the results plotted in Figure 1. On the left side of the figure, the input data is challenging to categorize. However, when the Gaussian RBF (1) is applied to the data, categorization becomes trivially feasible through the use of a linear hyperplane.

Kullback-Leibler divergence

The Kullback-Leibler (KL) divergence serves as a measure of how different one continuous probability distribution is from another. Given two separate probability distributions $p(x)$ and $q(x)$ over the random variable x , the KL divergence [7] is given by

$$KL(p||q) = \mathbb{E}_{x \sim p} \left[\log \frac{p(x)}{q(x)} \right] = \mathbb{E}_{x \sim p} [\log p(x) - \log q(x)].$$

A noteworthy property of the KL divergence is its non-negativity, being equal to zero if and only if $p(x)$ and $q(x)$ are the same distribution. In SVGD, the KL divergence serves as an 'error measure' between two probability densities, quantifying the discrepancy between the current approximation and the target distribution.

Kernel Stein Discrepancy KSD

The foundational theoretical framework for SVGD comes from the concept of Kernelized Stein Discrepancy (KSD) [4]. KSD combines the potency of Stein's Identity [8] with the principles of Reproducing Kernel Hilbert Space (RKHS) theory [5]. This results in a computationally tractable measure of discrepancy between a ground truth distribution and samples from an arbitrary distribution.

Stein Identity

Consider a smooth density function $p(x)$ defined on the space X , and let $\phi(x) = [\phi_1(x), \phi_2(x), \dots, \phi_d(x)]^T$ be a smooth vector function. Stein's Identity [8] states that for a smooth density $p(x)$, the expected value of the Stein operator \mathcal{A}_p acting on $\phi(x)$ is zero:

$$\mathbb{E}_p [\mathcal{A}_p \phi(x)] = 0 \quad \text{where} \quad \mathcal{A}_p \phi(x) := \phi(x) \nabla_x \log p(x)^T + \nabla_x \phi(x).$$

Here we call \mathcal{A} as the *Stein Operator*, which is a function operator. Notably, \mathcal{A}_p can often be computed in practice, even for models with intractable normalization constants, owing to its dependency only on p through the score function $\nabla_x \log p(x)$. The linearity of the Stein operator, expressed as $\mathcal{A}(f + g) = \mathcal{A}f + \mathcal{A}g$, is a key property.

Stein Discrepancy

Now, if we consider the expedition is taken with respect to q , while \mathcal{A}_p is still related to p , then there must exist a function ϕ such that it holds

$$\mathbb{E}_{x \sim q} [\mathcal{A}_p \phi(x)] = \mathbb{E}_{x \sim q} [\mathcal{A}_p \phi(x)] - \mathbb{E}_{x \sim q} [\mathcal{A}_q \phi(x)] \quad (2)$$

$$= \mathbb{E}_{x \sim q} [(\nabla_x \log p(x) - \nabla_x \log q(x)) \phi(x)^T]. \quad (3)$$

Equation (3) represents the differences between the score functions, implying non-zero values if and only if $p = q$. Moreover, it can be demonstrated [9] that this leads to

$$\mathbb{E}_{x \sim q} [\mathcal{A}_p \phi(x)] = \mathbb{E}_{x \sim q} [\text{trace}(\mathcal{A}_p \phi(x))]$$

where the trace is employed to form a scalar, though other suitable matrix norms can be used. This leads to the definition of the Stein Discrepancy measure:

$$S(q, p) = \max_{\phi \in \mathcal{H}^d, \|\phi\| \leq 1} \left\{ \mathbb{E}_p [\text{trace}(\mathcal{A}_p \phi(x))]^2 \right\}. \quad (4)$$

Here, we apply the condition that \mathcal{H}^d is a RKHS for some kernel k . It has been shown [4] that the optimal solution for Equation (4) is given by

$$\phi_{q,p}^*(\cdot) = \mathbb{E}_{x \sim q} [\mathcal{A}_p k(x, \cdot)]. \quad (5)$$

This result forms the cornerstone of the SVGD algorithm, empowering it with the ability to efficiently estimate the discrepancy between probability distributions.

Stein Variation Gradient Descent

It has been shown [10] that under a deterministic transition $z = T(x) + \varepsilon \phi(x)$ with $x \sim q(x)$ and using the ideas from Equation (2) can lead to

$$\nabla_\varepsilon \text{KL}(q_{[T]} || p) \Big|_{\varepsilon=0} = -\mathbb{E}_{x \sim q} [\text{trace}(\mathcal{A}_p \phi(x))]. \quad (6)$$

By utilising the change of variable formula for probability distribution, the distribution of $q_{[T]}$ can be expressed as

$$q_{[T]}(z) = q(T^{-1}(z)) \cdot \det |\nabla_z T^{-1}(z)|.$$

Using the results and concepts discussed in the previous sections, it can be shown that Equation (5) represents the optimal direction of perturbation in the unit ball of the RKHS \mathcal{H}^d . This direction minimizes the KL divergence of the transformed distribution $q_{[T]}$ optimally, with the magnitude change given by:

$$\nabla_\varepsilon \text{KL}(q_{[T]} || p) = -S(q, p) \quad (7)$$

In practice, applying this identity perturbation at every time step ε reduces the KL divergence by a factor of $\varepsilon S(q, p)$. Iterating this process for a sufficient number of steps leads to convergence to the true distribution p . This process can be expressed as an Ordinary Differential Equation:

$$\partial_t x = \hat{\phi}_{q,p}^*(x)$$

where $\hat{\phi}_{q,p}^*(x)$, representing the expectation as stated in Equation (5), can be empirically estimated by taking the mean of n sampled particles:

$$\hat{\phi}_{q,p}^*(x) = \frac{1}{n} \sum_{j=1}^n \left[k(x_j, x) \nabla_{x_j} \log p(x_j) + \nabla_{x_j} k(x_j, x) \right].$$

Hence we can construct an iterative update for each particle x_i by the scheme define in [10]

$$x_i \leftarrow x_i + \frac{\varepsilon}{n} \sum_{j=1}^n \left[k(x_j, x_i) \nabla_{x_j} \log p(x_j) + \nabla_{x_j} k(x_j, x_i) \right] \quad (8)$$

which is implemented in the Algorithm 1, which again is extracted from [10]. Furthermore, the term $k(x_j, x) \nabla_{x_j} \log p(x_j)$ can be interpreted as the force that drives the particles towards the regions of high probability in of the $p(x)$ distribution. On the other hand, the term $\nabla_{x_j} k(x_j, x)$ acts as a repulsive force between x_j and x_i which push the particles away from each other, ensuring the particles do not bunch together and sample the distribution more evenly. If we consider the case for a single particle $n = 1$ then Equation (8) reduces to the standard gradient descent for maximizing $\log p(x)$, which is referred to as the MAP. This is due to $\nabla_x k(x, x) = 0$, which holds for RBF (1). With the addition of more particles, SVGD interpolates between gradient descent and approximate inference. Two different methods have been presented to analyze SVGD theoretically. The first treats SVGD as the gradient flow of the KL divergence function in the space of probability measures metrized by a RKHS variation of the Wasserstein distance [11]. Alternatively, in [12], fixed points of SVGD are viewed as a type of quadratic quadrature method, from which an approximation bound can be derived.

Numerical Experiments

A simplified version of Algorithm 1 was implemented (see Code 2). The kernel was set to the Radial Basis Function (RBF) as in (1) with a bandwidth $h = 1$ for simplicity. The step size $\varepsilon = 0.01$ was chosen for a balance between accuracy and convergence speed. In [10], AdaGrad was used, and in [13], a coin betting algorithm determined an optimal step size. A toy example (see code 3) was set up with the target distribution $p(x) = \mathcal{N}(x - 4, 1/2) + \mathcal{N}(x; 2, 4)$, a bimodal Gaussian distribution. The iteration began with 100 particle samples from $\mathcal{N}(0, 1)$, visible in the histogram plot in Figure 2. Subsequent iterations (Figures 3 and 4) showcase convergence to the bimodal distribution of $p(x)$. Another example (see code 4) set the target as a two-dimensional Gaussian $p(x) = \mathcal{N}((2, 0)^T, ((2, 0), (0, 1))^T)$. Taking 100 particle samples from $\mathcal{N}((0, 0)^T, ((1, 0), (0, 1))^T)$ resulted in the plots shown in Figure 5. By removing $k(x_j, x) \nabla_{x_j} \log p(x_j)$, the particles were forced away (Figure 6). Similarly, removing $\nabla_{x_j} k(x_j, x)$ resulted in bunching of the particles (Figure 7).

References

1. Christian P Robert, George Casella, and George Casella. *Monte Carlo statistical methods*, volume 2. Springer, 1999.
2. Andrew Gelman, John B Carlin, Hal S Stern, and Donald B Rubin. Bayesian data analysis second edition corrected version (30 jan 2008). 1995.
3. Danilo Jimenez Rezende and Shakir Mohamed. Variational inference with normalizing flows, 2016.
4. Qiang Liu, Jason D. Lee, and Michael I. Jordan. A kernelized stein discrepancy for goodness-of-fit tests and model evaluation, 2016.
5. Alain Berlinet and Christine Thomas-Agnan. *Reproducing kernel Hilbert spaces in probability and statistics*. Springer Science & Business Media, 2011.
6. Erwin Kreyszig. *Introductory functional analysis with applications*, volume 17. John Wiley & Sons, 1991.
7. Ian Goodfellow et al. *Deep Learning*. MIT press, 2016.
8. Charles Stein. A bound for the error in the normal approximation to the distribution of a sum of dependent random variables. In *Proceedings of the Sixth Berkeley Symposium on Mathematical Statistics and Probability, Volume 2: Probability Theory*, volume 6, pages 583–603. University of California Press, 1972.
9. Jackson Gorham and Lester Mackey. Measuring sample quality with stein’s method, 2019.
10. Qiang Liu and Dilin Wang. Stein variational gradient descent: A general purpose bayesian inference algorithm, 2019.
11. Qiang Liu. Stein variational gradient descent as gradient flow, 2017.
12. Qiang Liu and Dilin Wang. Stein variational gradient descent as moment matching, 2018.
13. Louis Sharrock and Christopher Nemeth. Coin sampling: Gradient-based bayesian inference without learning rates, 2023.




Figure 1: Plots from code 1, on the left we have the two two-dimensional input data which can not be categorized linearly. Two on the right demonstrated that after applying a suitable kernel function 1 the data can be categorized linearly in a higher dimension.

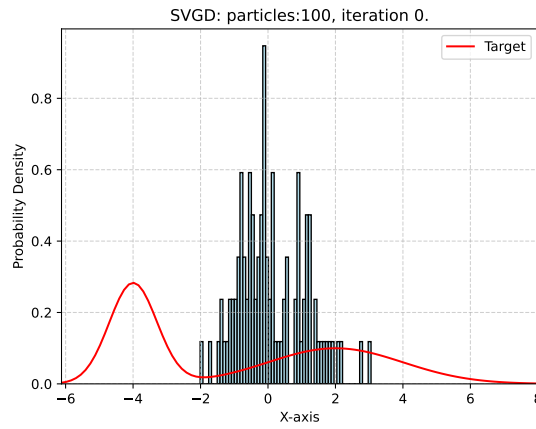


Figure 2: The histogram of the current sample of 100 particles with the red line being the target distribution. Here at iteration 0 we see that the histogram forms the initial sample distribution. From code 3.




Figure 3: The histogram of the current sample of 100 particles with the red line being the target distribution.




Figure 4: The histogram of the current sample of 100 particles with the red line being the target distribution.




Figure 5: Here we overlay the 100 particles on the two-dimensional plot of a Gaussian after 400 iteration of our SVGD. From code 4.

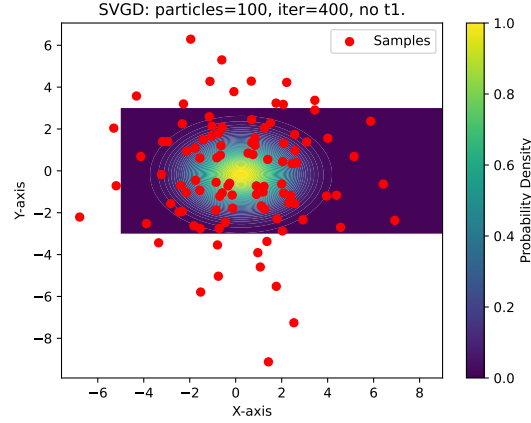


Figure 6: Removing the term that drives the particles towards the higher probability regions, we observed the spreading out as predicted due to the repulsive force between each particle. Here we overlay the 100 particles on the two-dimensional plot of a Gaussian after 400 iterations of our SVGD. From code 4.

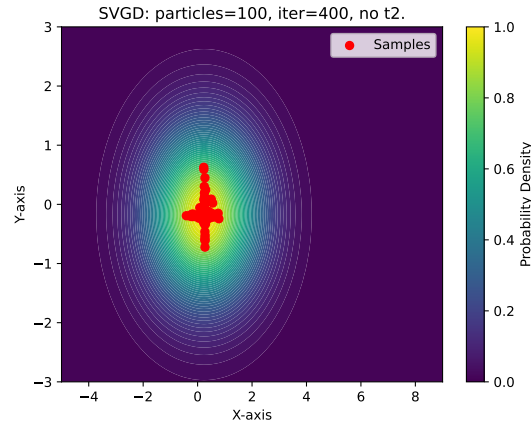


Figure 7: Removing the term that repels each particle from each other leads to the bunching of particles at the higher probability region. Here we overlay the 100 particles on the two-dimensional plot of a Gaussian after 400 iterations of our SVGD. From code 4.

132 B Algorithm

Algorithm 1 Bayesian Inference via Variational Gradient Descent. [10]

Require: A target distribution with density function $p(x)$ and a set of initial particles $\{x_i^0\}_{i=1}^n$

Ensure: A set of particles $\{x_i\}_{i=1}^n$ that approximates the target distribution.

```

1: for iteration = 1 to  $l$  do
2:   for  $i = 1$  to  $n$  do
3:      $x'_i \leftarrow x_i + \epsilon \hat{\phi}^*(x_i)$  ▷ Update particles
4:     where  $\hat{\phi}^*(x) = \frac{1}{n} \sum_{j=1}^n [k(x'_j, x) \nabla_{x_j} \log p(x_j) + \nabla_{x_j} k(x'_j, x)]$ 
5:   end for
6: end for

```

133 C Python Code

Listing 1: Kernel example.

```

134 # Generate synthetic data
135 X, y = make_circles(n_samples=100, factor=.1, noise=.1)
136 Z = feature_map_2(X)
137
138 # 2D scatter plot
139 fig1 = plt.figure(figsize=(8, 8))
140 ax1 = fig1.add_subplot(1, 1, 1)
141 ax1.scatter(X[:, 0], X[:, 1], c=y, cmap='viridis')
142 ax1.set_xlabel('$x_1$')
143 ax1.set_ylabel('$x_2$')
144 ax1.set_title('Original_dataset')
145
146 # Show the first plot
147 plt.show()
148
149 # 3D scatter plot
150 fig2 = plt.figure(figsize=(16, 8))
151 ax2 = fig2.add_subplot(1, 2, 1, projection='3d')
152 ax2.scatter3D(Z[:, 0], Z[:, 1], Z[:, 2], c=y, cmap='viridis')
153 ax2.set_xlabel('$z_1$')
154 ax2.set_ylabel('$z_2$')
155 ax2.set_zlabel('$z_3$')
156 ax2.set_title('Transformed_dataset')
157
158 # Draw a hyperplane at z
159 z_hyperplane = 0.65
160 xx, yy = np.meshgrid(np.linspace(min(Z[:, 0]), max(Z[:, 0]), 50),
161                      np.linspace(min(Z[:, 1]), max(Z[:, 1]), 50))
162 zz = np.full_like(xx, z_hyperplane)
163 ax2.plot_surface(xx, yy, zz, alpha=0.5, color='red', label='Hyperplane')
164
165 # Show the second plot
166 plt.show()
167
168 # 3D scatter plot (side view)
169 fig3 = plt.figure(figsize=(16, 8))

```



```

170 ax3 = fig3.add_subplot(1, 3, 1, projection='3d')
171 ax3.scatter3D(Z[:, 0], Z[:, 1], Z[:, 2], c=y, cmap='viridis')
172
173 # Draw a hyperplane at z=0.5
174 ax3.plot_surface(xx, yy, zz, alpha=0.5, color='red', label='Hyperplane')
175
176 ax3.set_xlabel('$z_1$')
177 ax3.set_ylabel('$z_2$')
178 ax3.set_zlabel('$z_3$')
179 ax3.set_title('Transformed_dataset_(Side_View)')
180
181 # Set the view to the side
182 ax3.view_init(elev=0, azimuth=90)
183
184
185
186 # Adjust layout for better spacing
187 plt.tight_layout()
188 plt.show()

```

Listing 2: Simple SVGD implemented

```

189 import numpy as np
190 import matplotlib.pyplot as plt
191 import numpy.matlib as nm
192 def rbf_kernel(x, y, h=1):
193     # Radial basis function (RBF) kernel
194     return np.exp(-np.sum((x - y)**2) / (2 * h**2))
195
196 def compute_Kxy(x, h=1):
197     # K xy
198     # Dist function for all particles
199     # dim(x) = N then we have a ( N x N) size matrix, with ones on the diag.
200
201     N = len(x)
202     Kxy = np.zeros((N, N))
203
204     for i in range(N):
205         for j in range(i): # Only compute the lower half
206             Kxy[i, j] = rbf_kernel(x[i], x[j], h)
207
208     # Fill in the upper half using symmetry
209     Kxy = Kxy + Kxy.T + np.diag(np.ones(N))
210
211     return Kxy
212
213 def grad_kxy(x, Kxy, h=1):
214     dx_kxy = -np.matmul(Kxy, x)
215     sumkxy = np.sum(Kxy, axis=1)
216
217     for i in range(x.shape[1]):
218         dx_kxy[:, i] = dx_kxy[:, i] + ( x[:, i] * sumkxy )

```

```

219     #dx_kxy /= h**2
220
221     return dx_kxy
222
223 def svgd_grad_X(x, log_prob_func,):
224     x = np.copy(x)
225     # Computes  $\phi^*$  for each time iteration.
226
227     log_prob_grad = log_prob_func(x)
228
229     # Compute the kernel matrix
230     Kxy = compute_Kxy(x)
231     dx_Kxy = grad_kxy(x, Kxy)
232
233     term1 = np.matmul(Kxy, log_prob_grad)
234     term2 = dx_Kxy
235
236     grad_x = (( term1 + term2) / x.shape[0])
237
238     return grad_x
239
240
241 def svgd_run(x0, d_log_prob_func, n_iter=1, step_size=1e-3, alpha = 1):
242
243     # adagrad with momentum
244     fudge_factor = 1e-6
245
246     x = np.copy(x0)
247     x_hist = [np.ones_like(x)]*n_iter
248     historical_grad =0
249     #
250     for n in range(n_iter):
251         x_hist[n] = x
252         grad_x = svgd_grad_X(x, d_log_prob_func)
253
254         if n ==0:
255             historical_grad = historical_grad + grad_x**2
256         else:
257             historical_grad = alpha * historical_grad + (1 - alpha) * (grad_x **
258
259         # force
260         F = np.divide(grad_x, fudge_factor+np.sqrt(historical_grad))
261
262         # update
263         x = x + step_size * F
264
265     return x, x_hist

```

Listing 3: Bimodal Gaussian Example

```

266 # Bimodal
267 mul = -4

```

```

268 var1 = 0.5
269 mu2 = 2
270 var2 = 4
271
272 def bimodal_dlnprob(x):
273     return (-1*(x-mu1)*(1/var1)*bimodal_gaussian(x, mu1, var1) - \
274            (x-mu2)*(1/var2)*bimodal_gaussian(x, mu2, var2)) / \
275            (bimodal_gaussian(x, mu1, var1) + bimodal_gaussian(x, mu2, var2))
276
277 x0 = np.random.normal(0,1, [100,1])
278
279 x, xh=svgd_run(x0, bimodal_dlnprob,500,0.01)
280
281 print("svgd: _mu=_{} _var=_{}".format(round(np.mean(x),2), round(np.std(x)**2,2))
282 plot_results(mu1, var1, mu2, var2, x)

```

Listing 4: 2D Unimodal Gaussian

```

283 # Unimodal Gaussian
284 A = np.array([[2,0],[0,1]])
285 mu = np.array([2,0])
286
287 def unimodal_gaussian_d_log_prob(x):
288     return -1*np.matmul(x-nm.repmat(mu, x.shape[0], 1), np.linalg.inv(A))
289
290 np.random.seed(10)
291 x0 = np.random.normal(0,1, [50,2]);
292 xx, xh=svgd_run(x0, unimodal_gaussian_d_log_prob,500,0.01)

```